

# **EE7204: COMPUTER VISION AND IMAGE PROCESSING**

**TAKE HOME ASSIGNMENT - 01**

NAME : LAKSHAN D.P.W.K.

REG. NO : EG/2019/3649

SEMESTER : 07

DATE : 15/03/2024

# Contents

<b>List of Figures.....</b>	<b>3</b>
<b>1     GitHub Link .....</b>	<b>4</b>
<b>2     Coding Answers .....</b>	<b>4</b>
2.1 To reduce the number of intensity levels in an image from 256 to 2, in integer powers of 2. The desired number of intensity levels needs to be a variable input to your program....	4
2.2 Load an image and then perform a simple spatial 3x3 average of image pixels. Repeat the process for a 10x10 neighborhood and again for a 20x20 neighborhood. ....	6
2.3 Rotate an image by 45 and 90 degrees.....	7
2.4 For every 3×3 block of the image (without overlapping), replace all corresponding 9 pixels by their average. This operation simulates reducing the image spatial resolution. Repeat this for 5×5 blocks and 7×7 blocks. ....	8
<b>3     Results (Captured Images) .....</b>	<b>10</b>
3.1 To reduce the number of intensity levels in an image from 256 to 2, in integer powers of 2. The desired number of intensity levels needs to be a variable input to your program..	10
3.1.1 Gray-Scale Results.....	10
3.1.2 Color Results.....	14
3.2 Load an image and then perform a simple spatial 3x3 average of image pixels. Repeat the process for a 10x10 neighborhood and again for a 20x20 neighborhood. ....	19
3.3 Rotate an image by 45 and 90 degrees.....	21
3.4 For every 3×3 block of the image (without overlapping), replace all corresponding 9 pixels by their average. This operation simulates reducing the image spatial resolution. Repeat this for 5×5 blocks and 7×7 blocks. ....	24

## List of Figures

FIGURE 2.1: ORIGINAL GRAYSCALE IMAGE .....	10
FIGURE 2.2: 256 LEVELS REDUCED INTENSITY IMAGE.....	10
FIGURE 2.3: 128 LEVELS REDUCED INTENSITY IMAGE.....	11
FIGURE 2.4: 64 LEVELS REDUCED INTENSITY IMAGE.....	11
FIGURE 2.5: 32 LEVELS REDUCED INTENSITY IMAGE.....	12
FIGURE 2.6: 16 LEVELS REDUCED INTENSITY IMAGE.....	12
FIGURE 2.7: 8 LEVELS REDUCED INTENSITY IMAGE.....	13
FIGURE 2.8: 4 LEVELS REDUCED INTENSITY IMAGE.....	13
FIGURE 2.9: 2 LEVELS REDUCED INTENSITY IMAGE.....	14
FIGURE 2.10: ORIGINAL IMAGE.....	14
FIGURE 2.11: 256 LEVELS REDUCED INTENSITY IMAGE.....	15
FIGURE 2.12: 128 LEVELS REDUCED INTENSITY IMAGE.....	15
FIGURE 2.13: 64 LEVELS REDUCED INTENSITY IMAGE.....	16
FIGURE 2.14: 32 LEVELS REDUCED INTENSITY IMAGE.....	16
FIGURE 2.15: 16 LEVELS REDUCED INTENSITY IMAGE.....	17
FIGURE 2.16: 8 LEVELS REDUCED INTENSITY IMAGE.....	17
FIGURE 2.17: 4 LEVELS REDUCED INTENSITY IMAGE.....	18
FIGURE 2.18: 2 LEVELS REDUCED INTENSITY IMAGE.....	18
FIGURE 2.19: ORIGINAL IMAGE (IN GRAYSCALE) .....	19
FIGURE 2.20: 3 X 3 SPATIAL AVERAGE IMAGE .....	19
FIGURE 2.21: 10 X 10 SPATIAL AVERAGE IMAGE .....	20
FIGURE 2.22: 20 X 20 SPATIAL AVERAGE IMAGE .....	20
FIGURE 2.23: ORIGINAL IMAGE.....	21
FIGURE 2.24: IMAGE ROTATED BY 45 DEGREES CLOCKWISE .....	21
FIGURE 2.25: IMAGE ROTATED BY 45 DEGREES COUNTERCLOCKWISE.....	22
FIGURE 2.26: IMAGE ROTATED BY 90 DEGREES CLOCKWISE .....	22
FIGURE 2.27: IMAGE ROTATED BY 90 DEGREES COUNTERCLOCKWISE.....	23
FIGURE 2.28: ORIGINAL IMAGE (IN GRAYSCALE) .....	24
FIGURE 2.29: 3 X 3 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING).....	24
FIGURE 2.30: 5 X 5 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING).....	25
FIGURE 2.31: 7 X 7 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING).....	25

# 1 GitHub Link

[https://github.com/KalinduLakshan/Assignment\\_01.git](https://github.com/KalinduLakshan/Assignment_01.git)

## 2 Coding Answers

- 2.1 To reduce the number of intensity levels in an image from 256 to 2, in integer powers of 2. The desired number of intensity levels needs to be a variable input to your program.**

```
import cv2
import numpy as np

# Define the reduce_intensity_levels function
def reduce_intensity_levels(image, q_levels):
    # Calculate the scaling factor
    scale = 255 / (q_levels - 1)

    # Quantize pixel values
    quantized_image = np.round(image / scale) * scale

    # Convert to uint8
    quantized_image = quantized_image.astype(np.uint8)

    return quantized_image

# Loading the image file, assigned the image file name to a variable for a
# gray scale or color image input
# identification
image_path = 'sample.jpg'

isGray = int(input("Do you need to convert the original image into a gray
scale image (Put 1='yes', 0='no'): "))

if isGray:
    input_image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)

    # Check whether the image can be read or not
    if input_image is None:
        print("Sorry: Your given image could not read.")
    else:
        # Enter the number of intensity levels as a power of 2
        Q_levels = int(input("Enter the number of intensity levels (power of
2): "))

        # Check whether the number of intensity levels is a power of 2 or not
        if Q_levels & (Q_levels - 1) != 0:
            print("Error: Number of intensity levels must be a power of 2.")
        else:
            # Reduce the intensity levels by changing the quantization levels
```

```

        output_image = reduce_intensity_levels(input_image, Q_levels)

        # Display the original and quantized images
        cv2.imshow('Original Image', input_image)
        cv2.imshow(f'{Q_levels} Levels Reduced Intensity Image',
output_image)
        cv2.waitKey(0)
        cv2.destroyAllWindows()

else:
    input_image = cv2.imread(image_path)

    # Check whether the input image can be read or not
    if input_image is None:
        print("Sorry: Your given image could not read.")
    else:

        # Enter the number of intensity levels as a power of 2
        Q_levels = int(input("Enter the number of intensity levels (power of
2): "))

        # Check whether the number of intensity levels is a power of 2 or not
        if Q_levels & (Q_levels - 1) != 0:
            print("Error: Number of intensity levels must be a power of 2.")
        else:
            # Reduce the intensity levels by changing the quantization levels
            output_image = reduce_intensity_levels(input_image, Q_levels)

            # Display the original and quantized images
            cv2.imshow('Original Image', input_image)
            cv2.imshow(f'{Q_levels} Levels Reduced Intensity Image',
output_image)

            cv2.waitKey(0)
            cv2.destroyAllWindows()

```

## 2.2 Load an image and then perform a simple spatial 3x3 average of image pixels. Repeat the process for a 10x10 neighborhood and again for a 20x20 neighborhood.

```
import cv2
import numpy as np

# Define the simple_spatial_average function
def simple_spatial_average(image, neighborhood_size):
    # Pad the image to handle the borders of the image
    padded_image = cv2.copyMakeBorder(image, neighborhood_size//2,
    neighborhood_size//2, neighborhood_size//2, neighborhood_size//2,
    cv2.BORDER_CONSTANT)

    # Initialize the output image
    output_image = np.zeros_like(image)

    # Apply the spatial averaging for pixels
    for y in range(image.shape[0]):
        for x in range(image.shape[1]):
            # Extract the current neighborhood
            neighborhood = padded_image[y:y+neighborhood_size,
            x:x+neighborhood_size]

            # Calculate average and assign it to the output pixel
            output_image[y, x] = np.mean(neighborhood)

    return output_image

# Load the input image
image_path = 'sample.jpg'
input_image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)

if input_image is None:
    print("Sorry: Your given image could not read.")
else:
    cv2.imshow('Original Image', input_image)
    # Perform spatial averaging for 3, 10, 20 neighborhood sizes
    neighborhood_sizes = [3, 10, 20]

    # Loop through different neighborhood sizes
    for size in neighborhood_sizes:
        output_image = simple_spatial_average(input_image, size)

        # Display the result
        cv2.imshow(f'{size}x{size} Neighborhood Image', output_image)
    cv2.waitKey(0)
    cv2.destroyAllWindows()
```

## 2.3 Rotate an image by 45 and 90 degrees.

```
import cv2

# Loading the image file
input_file_name = 'cat.jpg'
input_image = cv2.imread(input_file_name)

# Check whether the input image can be read or not
if input_image is None:
    print("Sorry: Your given image could not read.")
else:
    # Original Image
    cv2.imshow("Original Image", input_image)

    # Rotate the image by 90 degrees clockwise
    output_90_img = cv2.rotate(input_image, cv2.ROTATE_90_CLOCKWISE)
    cv2.imshow("Image rotated by 90 clockwise", output_90_img)

    # Rotate the image by 90 degrees counterclockwise or 270 clockwise
    output_counter90_img = cv2.rotate(input_image,
cv2.ROTATE_90_COUNTERCLOCKWISE)
    cv2.imshow("Image rotated by 90 counterclockwise", output_counter90_img)

'''
45 degrees rotation(arbitrary rotation) cannot be performed using
cv2.rotate() method. Here we need to perform the general rotation by giving
the angle and the scale factor. Need to create a rotation matrix using
cv2.getRotationMatrix2D() and apply it to image with cv2.warpAffine()
'''

    # Rotate the image by 45 degrees clockwise
    (height, width, channel) = input_image.shape
    center = (width/2, height/2)

    clockwise_angle = 45
    scale = 1.0

    img_clockwise_matrix = cv2.getRotationMatrix2D(center, clockwise_angle,
scale)
    output_45_img = cv2.warpAffine(input_image, img_clockwise_matrix, (width,
height))
    cv2.imshow("Image rotated by 45 clockwise", output_45_img)

    # Rotate the image by 45 degrees clockwise
    counterclockwise_angle = -45

    img_counterclockwise_matrix = cv2.getRotationMatrix2D(center,
counterclockwise_angle, scale)
    output_counter_45_img = cv2.warpAffine(input_image,
img_counterclockwise_matrix, (width, height))
    cv2.imshow("Image rotated by 45 counterclockwise", output_counter_45_img)

cv2.waitKey(0)
cv2.destroyAllWindows()
```

**2.4 For every 3×3 block of the image (without overlapping), replace all corresponding 9 pixels by their average. This operation simulates reducing the image spatial resolution. Repeat this for 5×5 blocks and 7×7 blocks.**

```
import cv2
import numpy as np

# Define the spatial_resolution_reduction
def spatial_resolution_reduction(image, block_size):
    # Get the dimensions of the input image
    height, width = image.shape[:2]

    # Calculate the number of blocks in both dimensions
    num_blocks_x = width // block_size # Integer value is taken
    num_blocks_y = height // block_size

    # Initialize the output image
    output_res_reduction_image = np.zeros_like(image)

    # Iterate over each block
    for y in range(num_blocks_y):
        for x in range(num_blocks_x):
            # Define the coordinates of the current block
            block_top = y * block_size
            block_bottom = (y + 1) * block_size
            block_left = x * block_size
            block_right = (x + 1) * block_size

            # Extract the current block from the image
            current_block = image[block_top:block_bottom,
block_left:block_right]

            # Calculate the average pixel value of the block
            average_value = np.mean(current_block)

            # Assign the average value to all pixels in the current block
            output_res_reduction_image[block_top:block_bottom,
block_left:block_right] = average_value

    return output_res_reduction_image

'''
Here I have used grayscale for better understandings because while reading
images using OpenCV, some color values
are manipulated.
Therefore the original color image has converted into a grayscale image for
better understanding
'''
```



```
# Load the input image
image_name = 'cat.jpg' # Image path
input_image_read = cv2.imread(image_name)
input_image = cv2.cvtColor(input_image_read, cv2.COLOR_BGR2GRAY)

if input_image is None:
    print("Sorry: Your given image could not read.")
else:
    # Original Image
    cv2.imshow("Original Image", input_image)

    # Perform spatial resolution reduction for different block sizes
    block_sizes = [3, 5, 7]

    for size in block_sizes:
        output_image = spatial_resolution_reduction(input_image, size)

        # Display result
        cv2.imshow(f'{size}x{size} Block Size Image', output_image)

cv2.waitKey(0)
cv2.destroyAllWindows()
```

### 3 Results (Captured Images)

- 3.1 To reduce the number of intensity levels in an image from 256 to 2, in integer powers of 2. The desired number of intensity levels needs to be a variable input to your program.

#### 3.1.1 Gray-Scale Results




FIGURE 3.1: ORIGINAL GRAYSCALE IMAGE




FIGURE 3.2: 256 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.3: 128 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.4: 64 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.5: 32 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.6: 16 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.7: 8 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.8: 4 LEVELS REDUCED INTENSITY IMAGE





FIGURE 3.9: 2 LEVELS REDUCED INTENSITY IMAGE

### 3.1.2 Color Results




FIGURE 3.10: ORIGINAL IMAGE




FIGURE 3.11: 256 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.12: 128 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.13: 64 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.14: 32 LEVELS REDUCED INTENSITY IMAGE






FIGURE 3.15: 16 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.16: 8 LEVELS REDUCED INTENSITY IMAGE




FIGURE 3.17: 4 LEVELS REDUCED INTENSITY IMAGE



FIGURE 3.18: 2 LEVELS REDUCED INTENSITY IMAGE

**3.2 Load an image and then perform a simple spatial 3x3 average of image pixels. Repeat the process for a 10x10 neighborhood and again for a 20x20 neighborhood.**




FIGURE 3.19: ORIGINAL IMAGE (IN GRAYSCALE)




FIGURE 3.20: 3 X 3 SPATIAL AVERAGE IMAGE




FIGURE 3.21: 10 X 10 SPATIAL AVERAGE IMAGE




FIGURE 3.22: 20 X 20 SPATIAL AVERAGE IMAGE



### 3.3 Rotate an image by 45 and 90 degrees.




FIGURE 3.23: ORIGINAL IMAGE




FIGURE 3.24: IMAGE ROTATED BY 45 DEGREES CLOCKWISE




FIGURE 3.25: IMAGE ROTATED BY 45 DEGREES COUNTERCLOCKWISE



FIGURE 3.26: IMAGE ROTATED BY 90 DEGREES CLOCKWISE



FIGURE 3.27: IMAGE ROTATED BY 90 DEGREES COUNTERCLOCKWISE

- 3.4 For every  $3 \times 3$  block of the image (without overlapping), replace all corresponding 9 pixels by their average. This operation simulates reducing the image spatial resolution. Repeat this for  $5 \times 5$  blocks and  $7 \times 7$  blocks.**



FIGURE 3.28: ORIGINAL IMAGE (IN GRAYSCALE)



FIGURE 3.29: 3 X 3 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING)





FIGURE 3.30: 5 X 5 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING)



FIGURE 3.31: 7 X 7 BLOCK OF THE IMAGE (WITHOUT OVERLAPPING)